



LAYERED PARTITIONING

WINDOWS
TECHNOLOGY

Josip Šaban, Hypo Alpe Adria AG, Klagenfurt



About

- Finished FER in year 2004, Cotrugli MBA in year 2011
- Holds all certificates for SQL Server since version 2000
- Also holds .NET, Project Server and Biztalk certifications
- Certified PM on CompTIA and PRINCE2 methodologies
- Works as a senior BI analyst in Hypo-Alpe-Adria, Klagenfurt
- Also tries to help grow his own company in Zagreb
- Finishing doctoral studies on FOI and preparing PMP exam

Summary

Partitioning is the one of the **key methods** that can enhance query performance, but there is no guarantee - we show the why and how of partitioning in very large data sets.

Other key methods we need are **indexing** and **data compression**.

A real – world case

- Help! My table is getting too big!
- One installation has a reporting system containing both fact and dimension tables
 - Fact tables are 100+ GB in size (including all indexes) quickly approaching 1 billion rows
 - Dimension tables are less than 1 GB in size (including all indexes) with less than 100 thousand rows
- A middle tier application has been designed to dynamically create and execute queries to run reports custom-designed by clients

What are very large data sets?

- There is no official definition
- Typically occupying TB range, OLAP or OLTP with large amount of users
- Billions of rows

What is too big where tables are concerned?

- SQL Server can store Multi-TB tables, so the answer is always contextual
- Big enough to cause performance issues (on the *current* hardware)
- Big enough to cause maintenance operations to take too long or simply not be practical at all (on the *current* hardware)
- Big due to containing large amounts of historic / "processed" / "completed" data

What options do I have?

- Option 1: Do nothing (the told you so/job gamble option)
- Option 2: Upgrade the hardware (the sticky plaster option)
- Option 3: If no extra disk space is available (the bail out option):
 - Delete the older portion of the data
 - Move the older portion of the data to Windows SQL Database (Azure)
- Option 4: If disk space is available (the scale out option):
 - Partition the data
 - Distributed partitioned views
 - Partitioned views
 - Partitioning (Enterprise only)

Agenda

- Setting the stage
 - Introduction to indexing
 - Introduction to compression
- Main event...partitioning in very large data sets
 - Change the way of thinking and designing large tables
 - Introduction to partitioning
 - Advanced partitioning

Setting the stage

The story of indexing

INTRODUCTION TO INDEXING

Indexing basics

- Index Considerations
 - Can dramatically increase query performance
 - Adds overhead for index maintenance
- Best Practices
 - Base design on real-world workloads
 - Scenarios:
 - Retrieving ranges of data
 - Retrieving specific values

Basic indexes

- Clustered index
 - Controls the physical order of rows
 - Does not require disk space
 - One per table (may include multiple columns)
 - Created by default on tables' Primary Key column

Basic indexes

- Non-Clustered Index
 - Physical data structures that facilitate data retrieval
 - Can have many indexes
 - Indexes may include many columns
 - Covering index
 - Filtered index

Advanced indexes

- FULLTEXT index
 - Stores the info about significant words and their location within the columns of a database table
- SPATIAL Index
 - Provides the ability to perform certain operations more efficiently on spatial objects (spatial data or Geometry data type) in a column of the geometry data type

Advanced indexes

- XML Index
 - XML indexes can be created on xml data type columns
- Columnstore Index
 - New feature of SQL 2012
 - This index does not use the B-Tree structure but column-oriented storages
 - <https://msdn.microsoft.com/en-us/library/gg492088.aspx>

The story of compression

INTRODUCTION TO COMPRESSION

Basic compression

- Why compress
 - Disk throughput much slower than memory and CPU
 - Need less disk space
 - Two kinds of compression
 - Backup compression (not discussed here)
 - Data compression

Enabling compression

- Alter Table [TableName] Rebuild Partition = All with (Data_compression = Compression Type on Partitions (x to n))
 - Compression Types
 - Row
 - Page
 - None
- Alter Index [IndexName] on [TableName] Rebuild with (Data_compression = Compression Type)
 - Compression Types
 - Row
 - Page
 - None

Page compression – used in DWH

Pre-Fix

Dictionary

CI structure

Page Header		
aaabbb	aaaabb	abcd
aaabcc	0bbb	abcd
aaaacc	aaaacc	0bbb

Database Compression Cost/Benefit Analysis

Benefits	Cost
Performance improvements	Increased CPU utilisation
<ul style="list-style-type: none">• More data in memory	
<ul style="list-style-type: none">• Reduced I/O	
Reduced disk space usage	
<ul style="list-style-type: none">• Database data files	
<ul style="list-style-type: none">• Backup files	
Reduced time to backup	
Cost Savings	

How to choose

- `sp_estimate_data_compression_savings`
- Quick Rule of thumb :
 - ROW is low-cost, generally 10% CPU overhead. Use it on everything on OLTP
 - PAGE is more expensive, but compresses more. Use it on everything on DWH
- Compression strategy documentation
 - [https://technet.microsoft.com/en-us/library/dd894051\(v=sql.100\).aspx](https://technet.microsoft.com/en-us/library/dd894051(v=sql.100).aspx)

How to choose - effectiveness

- Good Compression
- Numeric or fixed length character fields that don't use all the allocated bytes
- Repeating data or prefix values
- Poor or no Compression
- Fields using up all the allocated bytes
- Not much repeated data
- Repeated with non-repeating prefixes
- Out of row data
- FILESTREAM data



The main event

Why and how of partitioning

INTRODUCTION TO PARTITIONING


What is partitioning?

- In SQL Server all tables have at least one partition
 - “In SQL Server, all tables and indexes in a database are considered partitioned, even if they are made up of only one partition. Essentially, partitions form the basic unit of organization in the physical architecture of tables and indexes. This means that the logical and physical architecture of tables and indexes comprised of multiple partitions mirrors that of single-partition tables and indexes.”
 - Partitioned Table and Index Concepts
 - <https://msdn.microsoft.com/en-us/library/ms190787.aspx>
 - Partitioned Table and Index Strategies Using SQL Server 2008
 - <http://msdn.microsoft.com/en-us/library/dd578580.aspx>

Advantages of partitioning

- Improves Query Performance
- Better Data Manageability and Cost Effective
- Deleting and Moving data is faster from partitions
- Narrow downs the index maintenance window
- OLTP/DSS – Operational/Non-Operational data

A standard table...



Partitioned views...

vSales (Option One)

```
SELECT * FROM Sales_2014  
UNION ALL  
SELECT * FROM Sales_2015
```

vSales (Option Two)

```
SELECT * FROM Sales  
WHERE Year = '2014'  
UNION ALL  
SELECT * FROM Sales  
WHERE Year = '2015'
```

First SELECT

Second SELECT

Sales data 2014

Row 200

Row 400

...

Row 600.000

Sales data 2015

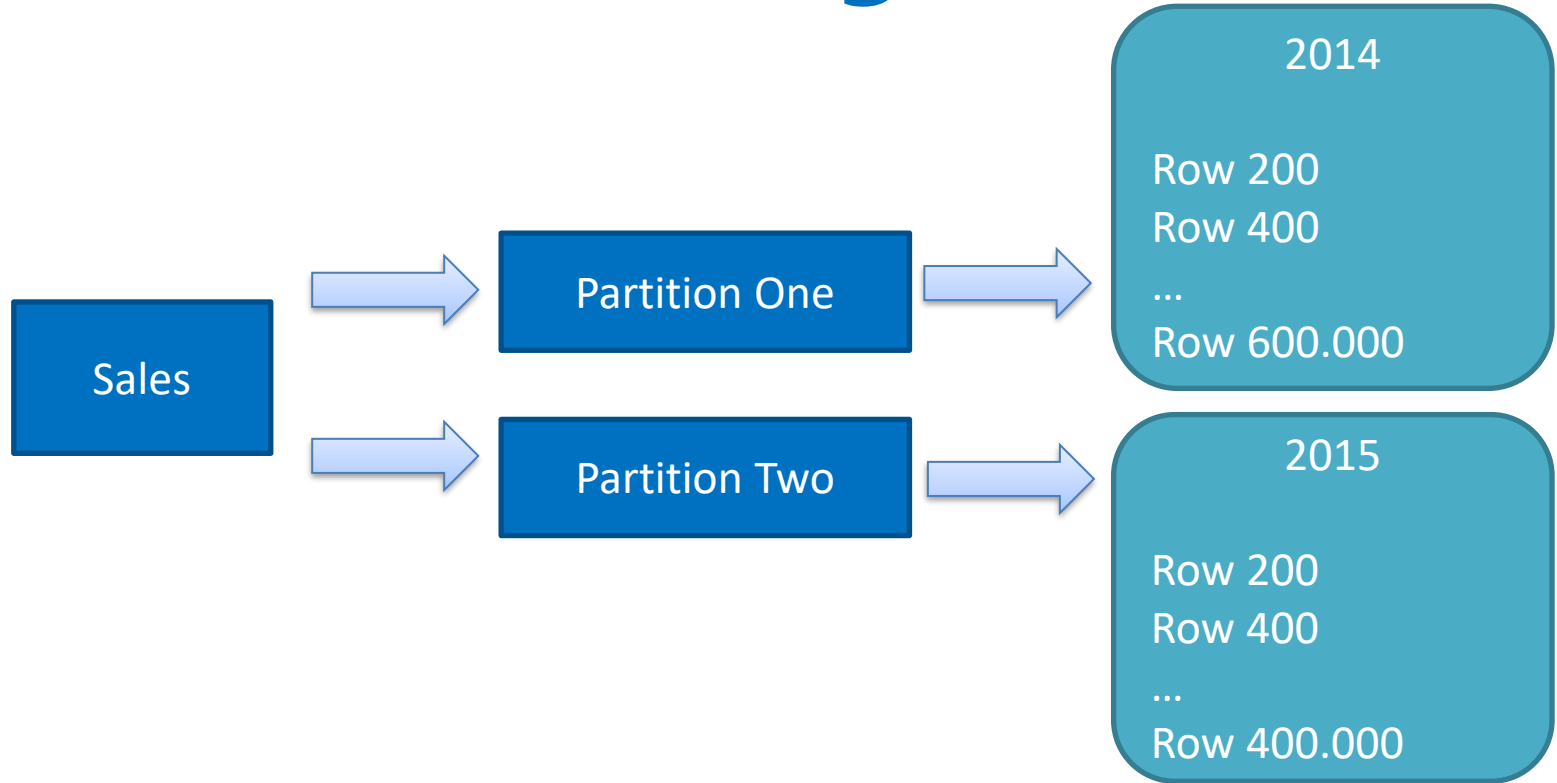
Row 200

Row 400

...

Row 400.000

Partitioning...



Partitioned views and Partitioning...

vSales (Option One)

```
SELECT * FROM Sales_2014  
UNION ALL  
SELECT * FROM Sales_2015
```

First SELECT

Sales data 2014

Row 200
Row 400
...
Row 600.000

January 2014

...

December 2014

vSales (Option Two)

```
SELECT * FROM Sales  
WHERE Year = '2014'  
UNION ALL  
SELECT * FROM Sales  
WHERE Year = '2015'
```

Second SELECT

Sales data 2015

Row 200
Row 400
...
Row 400.000

January 2015

...

December 2015

Partitioned Views Vs. Partitioning

- Partitioned Views
 - Any Version
 - Partition elimination
 - Different indexes per “partition”
 - Replication friendly (all types)
 - “Easy” feature
 - High IO during “partitioning”
 - Can support multiple constraints on multiple columns
- Partitioning
 - Enterprise only
 - Partition elimination
 - Same indexes for all partitions
 - Replication friendly (transactional)
 - “Complex” feature
 - Designed specifically to ease the pain of moving in and out large volumes of data
 - Partitioned on a single column

Data (Horizontal) partitioning

- You can partition ONLY by one column
- Large table/index can be split into multiple manageable portions
 - Horizontal partitioning takes groups of rows in a single table and allocates them in semi-independent physical sections
- SQL Server's horizontal partitioning is RANGE based

Horizontal ranges are based on a partition key

- A single column in the table
 - Use a computed column if you must, but make sure it performs well as a criterion and works for joins
- Typically a date or integer value
- Consider:
 - A column you will join on
 - A column you can always use as a criterion

Choosing a Partitioning Column

- Should reflect the best way to subdivide the target table and get a balanced distribution of data
- Used as a filter in most of the queries run against the table, otherwise you will not get the benefit of any partition elimination (accessing only the partitions needed as opposed to the whole table)
- Good candidates:
 - Date time columns: order date, inserted date, etc.
 - Countries: customer country

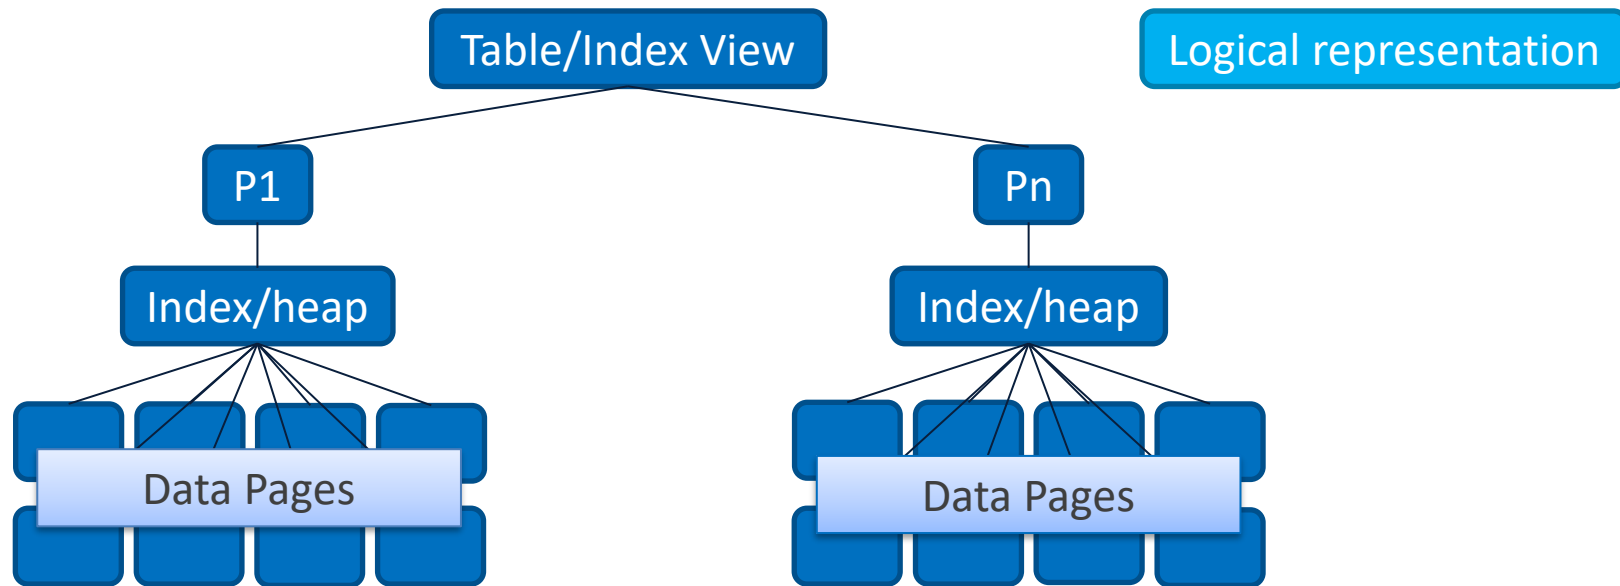
Partition concepts (1)

- PARTITION FUNCTION
 - Used to specify partition boundary values
 - Two types (represents boundary data directions)
 - LEFT – myself and my left range values
 - RIGHT – myself and my right range values
- PARTITION SCHEME
 - Logical & Partitions physically aligned
 - Span over single or multiple file groups
 - Can specify each partition can go to a individual file group or all partitions can go into a single file group

Partition concepts (2)

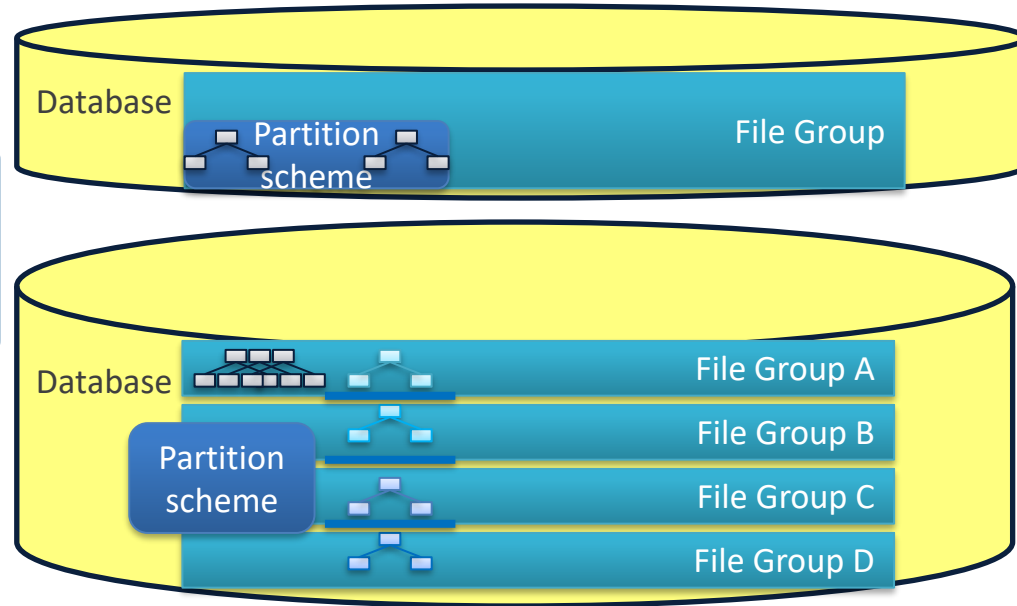
- SPLIT/MERGE
 - SPLIT
 - Introduces new boundary
 - Partition will be added to respective side (L/R)
 - MERGE
 - Deletes boundary
 - Partition will be merged to the respective side (L/R)
- SWITCH IN/OUT
 - Moving partition from partitioned table to other partitioned table called “in”
 - Moving partition from partitioned table to non partitioned table called “out”

Data Partitioning Architecture



Partition Physical Architecture


```
CREATE PARTITION SCHEME  
MyPartitionScheme_ps  
AS PARTITION MyPartitionFunctionName_pfn  
ALL TO ([FG]) -- specifying single file group  
-- TO ([FGA], [FGB]) -- specifying multiple file groups
```



Partitioning Functions


CREATE PARTITION FUNCTION

MyPartitionFunctionName_pfn(int) AS
RANGE LEFT FOR VALUES(10,50)




CREATE PARTITION FUNCTION


MyPartitionFunctionName_pfn(int) AS
RANGE RIGHT FOR VALUES(10,50)




SPLIT – LEFT boundary



```
ALTER PARTITION SCHEME MyPartitionScheme_ps  
NEXT USED [FG]; -- Specifying the group for new boundary  
ALTER PARTITION FUNCTION MyPartitionFunction_pfn()  
SPLIT RANGE(25)
```




MERGE – LEFT boundary



ALTER PARTITION FUNCTION

MyPartitionFunctionName_pfn()

MERGE RANGE (10)



Indexes can be created on the partition scheme...or not

Aligned Indexes

- Located on your partitioning scheme (or an identical partitioning scheme)
- Must contain the partitioning key
- If the partitioning key is not specified, it will be added for you. Note: this affects your primary key for the table!
- Indexes are aligned by default unless it is otherwise specified at creation time
- Perform better for aggregations and when partition elimination can be used

Non-aligned indexes

- Physically located elsewhere- either non partitioned or on a non-identical partitioning scheme
- May perform better with single-record lookup
- Allow unique indexes (because they do not have to contain the partitioning key)
- However, the presence of these preclude partition-switching!

Index rebuilds and compression


- Individual partitions *cannot* be rebuilt online
 - The entirety of a partitioned index *can* be rebuilt online
 - Individual partitions can be compressed
- For fact tables with archive data, older partitions can be rebuilt once with compression
 - Their file groups can then be made read-only

Switching

- Requires all indexes to be aligned
- Compatible with filtered indexes
- Data may be switched in or out only within the same file group
- Is a metadata-only operation requiring a schema modification lock
 - This can be blocked by DML operations, which require a schema stability lock
- Is an exceptionally fast way to load or remove a large amount of data from a table!
- <https://msdn.microsoft.com/en-us/library/ms191160.aspx>

SWITCH OUT

Partitioned table



Non Partitioned table

Partition 1

```
ALTER TABLE PartitionedTableName
```


```
SWITH PARTITION
```

```
2 TO NonPartitionedTableName
```


```
-- <Partition number> TO DestinationPartitionedTableName<destination Partition number>
```

SWITCH IN/OUT says “Hey, its your partition by updating metadata.”

Partitioned table



Previous Non Partitioned table



SQL Server Partitioning Myths (1)

- Myth 1: Partitioning is a "Scale-Out" solution
- Myth 2: Partitions must be created on different file groups
- Myth 3: To partition a non-partitioned table you will need to drop and recreate it
- Myth 4: Partitioning an existing table is a strictly offline operation

SQL Server Partitioning Myths (2)

- Myth 5: SWITCH-ing partitions OUT or IN in only a few seconds
- Myth 6: Altering a partition function is a metadata only operation
- Myth 7: Partitioned tables improve query performance
- Myth 8: Partitioned tables ease maintenance of VLDBs and Very Large Tables

Demo

- Let us start writing some code...
 - Simulating partitioning on Standard Edition
 - Partitioning demo
 - Concepts
 - Advanced demo

Partitioning without Enterprise Edition

- SQL Server 7.0 introduced partitioning through partitioned views
- In the world of partitioning, the number of rows queried is reduced according the granularity of the partitioning scheme
- There are, of course, some downsides
 - Queries that cross the partition rules will take longer
 - True partitioning requires SQL Server Enterprise



Contacts

- E-mail: josipsaban@gmail.com
- LinkedIn: www.linkedin.com/in/jsaban

